

# p3

---

Files included

## 1. Project Overview

### 1.1 Project Background and Objectives

### 1.2 Technology Stack

## 2. System Architecture Design

### 2.1 Overall Architecture

### 2.2 Database Architecture

## 3. Core Module Implementation

### 3.1 Database Design and Optimization

#### 3.1.1 Physical Model Design

#### 3.1.2 Index Optimization

#### 3.1.3 Partitioning Strategy

### 3.2 Dataset Construction

#### 3.2.1 Feature Generation

#### 3.2.2 Label Generation (Target Variable)

### 3.3 Model Training and Evaluation

#### 3.3.1 Feature Engineering for ML

#### 3.3.2 Model Selection and Training

#### 3.3.3 Model Evaluation

### 3.4 Cloud Deployment

#### 3.4.1 Aliyun OSS Setup

#### 3.4.2 Artifact Upload

#### 3.4.3 Deployment Workflow

## 4. System Testing and Optimization Results

### 4.1 Database Performance Optimization

### 4.2 Model Prediction Performance Analysis

## 5. Summary and Future Work

### 5.1 Project Achievements

## 5.2 Limitations and Future Directions

Name: Liu, Jiexi

Date: Nov,27,2025

NYU ID: N18245550

Couse section: Database Systems, Section 001

Project Part #3 Report

## Files included

- dataset/construct\_dataset.py
- dataset/customer\_insurance\_dataset.csv
- optimization/physical\_model.sql
- optimization/indexing.sql
- optimization/partitioning.sql
- prediction/model.pkl
- prediction/predict.py
- prediction/train\_model.py

## 1. Project Overview

### 1.1 Project Background and Objectives

The insurance industry is undergoing a significant digital transformation, with data-driven decision-making becoming a cornerstone of competitive advantage. Traditional methods of customer segmentation and insurance product recommendation are increasingly being augmented by advanced analytics and machine learning (ML) techniques. This project aims to build a comprehensive Customer Insurance Prediction System that leverages these technologies to predict customer insurance preferences and facilitate efficient business operations.

The core objectives of this project are:

- 1. Data Infrastructure:** Design, implement, and optimize a robust relational database to store and manage customer, account, and insurance policy data.
- 2. Data Acquisition & Preparation:** Construct a realistic, synthetic dataset that mirrors customer demographics and their corresponding insurance product holdings.
- 3. Machine Learning Model Development:** Train and evaluate a machine learning model to predict the type of insurance a customer is likely to purchase (e.g., Life, A&H, FSA, or None).
- 4. Cloud Deployment:** Deploy the trained model and associated datasets to a cloud-based data lake to ensure scalability, accessibility, and support for future operationalization.

## 1.2 Technology Stack

- **Database Management System:** PostgreSQL 14 (Local Deployment)
- **Programming Language:** Python 3.9
- **Data Processing & Analysis:** Pandas, NumPy
- **Synthetic Data Generation:** Faker library
- **Machine Learning:** Scikit-learn, XGBoost
- **Cloud Platform:** Alibaba Cloud (Aliyun)
- **Cloud Storage:** Alibaba Cloud Object Storage Service (OSS)
- **Development Tools:** VS Code, Navicat (for database management)

## 2. System Architecture Design

### 2.1 Overall Architecture

The system follows a layered, modular architecture to ensure separation of concerns, scalability, and maintainability. The architecture include:

- **Data Layer:** Responsible for the persistent storage of all structured data. A local PostgreSQL instance is used to maintain data integrity and support complex queries.
- **Processing Layer:** Encompasses data extraction, transformation, and loading (ETL) processes, as well as feature engineering to prepare data for machine learning.
- **Model Layer:** Where the machine learning model is trained, validated, and optimized.

This layer outputs a production-ready model artifact.

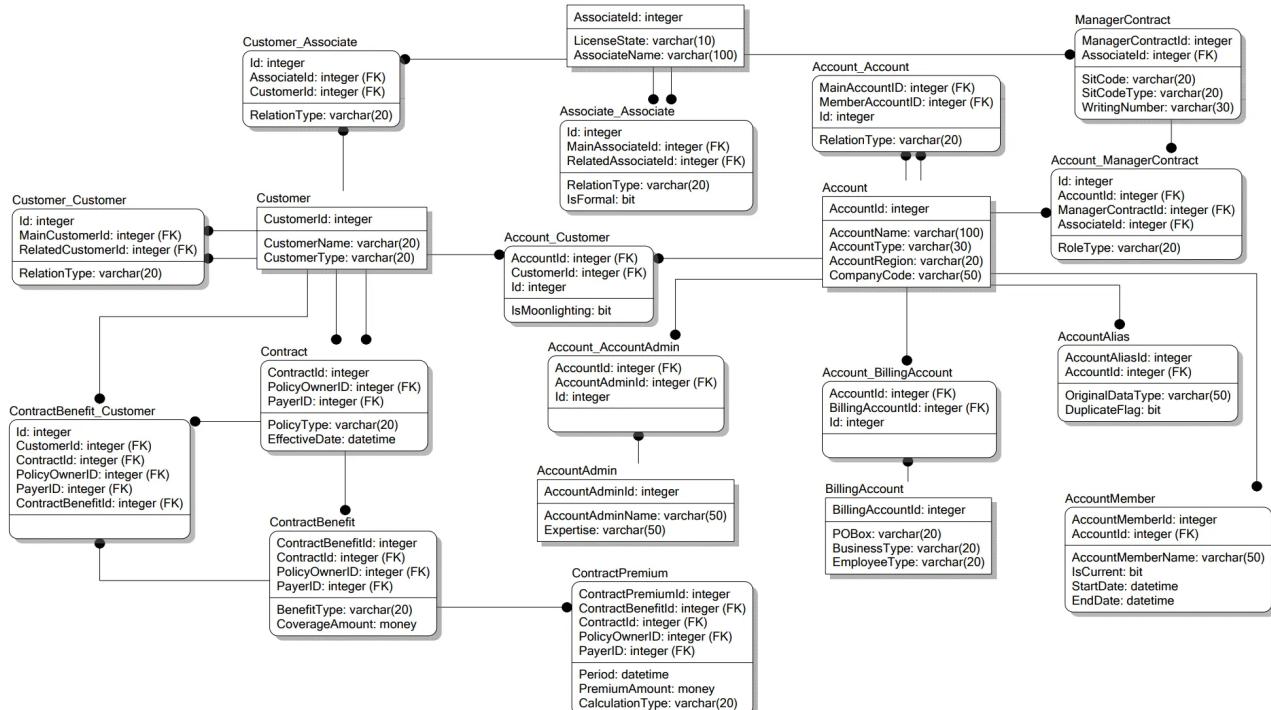
- **Deployment Layer:** Handles the transition of the model and data from the development environment to a production-ready cloud storage solution (Aliyun OSS), making them accessible for integration with business applications.

## 2.2 Database Architecture

The database schema is designed to accurately model the business entities and their relationships within an insurance context. It consists of 20+ normalized tables to minimize redundancy and ensure data consistency. Below is a detailed breakdown of the core tables:

- **Customer:** Stores demographic information about each customer (e.g., `CustomerId`, `Name`, `Age`, `Income`, `MaritalStatus`, `EducationLevel`).
- **Account:** Represents financial accounts associated with customers (e.g., `AccountId`, `CustomerId`, `AccountType`, `Region`, `Balance`).
- **Contract:** Records insurance policies held by customers (e.g., `ContractId`, `PolicyOwnerID`, `PayerID`, `EffectiveDate`, `PolicyType`, `PremiumAmount`).
- **Other Supporting Tables:** Include `Beneficiary`, `Payment`, `Claim`, etc., to capture the full breadth of insurance business operations.

The schema employs foreign key constraints extensively to enforce referential integrity (e.g., `Contract.PolicyOwnerID` references `Customer.CustomerId`).



### 3. Core Module Implementation

#### 3.1 Database Design and Optimization

##### 3.1.1 Physical Model Design

The physical database design translates the logical schema into efficient storage structures on disk. Key considerations included:

- **Data Types:** Choosing appropriate data types for each column to optimize storage and performance (e.g., `VARCHAR` for names, `INT` for ages, `DATE` for `EffectiveDate` ).
- **Constraints:**
  - **Primary Key (PK) Constraints:** Ensured unique identification of each record (e.g., `Customer.CustomerId` ).
  - **Foreign Key (FK) Constraints:** Maintained relational integrity between tables.
  - **Check Constraints:** Enforced business rules at the database level (e.g., `CHECK (Age BETWEEN 0 AND 120)` ).
  - **Not Null Constraints:** Ensured mandatory fields were populated.

##### 3.1.2 Index Optimization

To significantly improve the performance of read-heavy operations (queries), several strategic indexes were created:

- **Single-Column Indexes:** On frequently filtered columns like `Customer.Name` (`idx_customer_name`) and `Contract.PolicyType` (`idx_contract_policytype`).
- **Composite Indexes:** On columns frequently used together in `WHERE` clauses. For example, `idx_customer_type_age` on `(CustomerType, Age)` allows the database to quickly find all 'Individual' customers within a specific age range without scanning the entire table.
- **Join Indexes:** On foreign key columns to speed up `JOIN` operations between tables. For example, `idx_contract_policyowner_payer` on `(PolicyOwnerID, PayerID)` in the `Contract` table accelerates joins with the `Customer` table.

##### 3.1.3 Partitioning Strategy

For large tables expected to grow over time, partitioning was implemented to enhance query performance and manageability:

- **List Partitioning:** Applied to the `Account` table based on `AccountType` (e.g., 'Personal', 'Trust', 'Corporate'). Queries filtering by `AccountType` only scan the relevant partition.
- **Range Partitioning:** Applied to the `Contract` table based on `EffectiveDate`, creating monthly partitions for the years 2025 and 2026. This is particularly effective for time-series data, as reports or queries about a specific month only access that month's data.
- **Partitioned Indexes:** Local indexes were created on partitioned tables to ensure that index maintenance operations (e.g., `VACUUM`, `REINDEX`) could be performed on individual partitions, reducing lock contention and maintenance windows.

## 3.2 Dataset Construction

A synthetic yet realistic dataset was generated using Python's `Faker` library and custom logic embedded in `construct_dataset.py`. This dataset serves as the foundation for training the machine learning model.

### 3.2.1 Feature Generation

The dataset generation process created the following features for each customer:

- **Demographic Features:**
  - `Age` : Generated using a normal distribution (mean=45, std=15) clamped between 18 and 85.
  - `Gender` : Randomly assigned with a 50/50 split.
  - `Education_Level` : Categorical values ('High School', 'Bachelor's', 'Master's', 'PhD') assigned with weighted probabilities.
  - `Marital_Status` : Categorical values ('Single', 'Married', 'Divorced', 'Widowed') with age-dependent probabilities.
- **Socio-Economic Features:**
  - `Income` : A derived feature calculated as a base value plus bonuses based on `Age` (career progression) and `Education_Level` (higher education premium).
  - `Number_of_Dependents` : Derived from `Marital_Status` and `Age`, with married

customers having a higher chance of dependents.

- **Insurance Label:**

- `Insurance_Type` : The target variable indicating the customer's primary insurance product, generated using a rule-based probabilistic model.

### 3.2.2 Label Generation (Target Variable)

The `Insurance_Type` label (target variable) was generated using a sophisticated rule engine to mimic real-world insurance purchasing behavior:

1. **Base Weight Assignment:** Initial probabilities (weights) were assigned to each insurance type (e.g., Life, A&H, FSA, No\_Insurance).
2. **Rule-Based Adjustment:** Weights were dynamically adjusted based on customer features:
  - Older customers (>50) had increased weights for `Life` insurance.
  - Customers with dependents had increased weights for `Life` and `A&H` (Accident & Health).
  - High-income customers (> \$100K) had increased weights for `FSA` (Flexible Spending Account).
  - Young, single customers with no dependents had increased weights for `No_Insurance`.
3. **Probabilistic Sampling:** The final `Insurance_Type` for each customer was sampled from the adjusted probability distribution.
4. **Noise Injection:** A 3% random noise factor was introduced to ensure the dataset was not perfectly predictable and to reflect real-world variability.

The final dataset contained 10,000 customer records, each with features and a corresponding `Insurance_Type` label.

```
Successfully generated dataset: customer_insurance_dataset.csv
-----
Label Distribution:
Life: 38.1%
FSA: 27.5%
A&H: 19.7%
No Insurance: 14.7%
```

## 3.3 Model Training and Evaluation

### 3.3.1 Feature Engineering for ML

Before model training, the raw dataset underwent further preprocessing:

- **Feature Selection:** All generated features were retained based on domain knowledge.
- **Feature Transformation:**
  - **Categorical Features:** `Gender`, `Education_Level`, `Marital_Status` were one-hot encoded using `pandas.get_dummies()`.
  - **Numerical Features:** `Age`, `Income`, `Number_of_Dependents` were standardized (scaled to have mean=0 and std=1) using `sklearn.preprocessing.StandardScaler`.
- **Feature Creation:** A new feature `Income_Per_Dependent` was engineered as `Income / (Number_of_Dependents + 1)` to capture the financial burden per dependent.

### 3.3.2 Model Selection and Training

An XGBoost (Extreme Gradient Boosting) classifier was selected for this multi-class classification task due to its excellent performance, ability to handle non-linear relationships, and robustness to overfitting with proper tuning.

- **Model Initialization:** `XGBClassifier` from the `xgboost` library was initialized.
- **Hyperparameter Tuning:** Key hyperparameters were set based on initial experimentation and cross-validation:
  - `n_estimators=200` : Number of decision trees in the ensemble.
  - `learning_rate=0.05` : Step size shrinkage to prevent overfitting.
  - `max_depth=6` : Maximum depth of a tree, limiting its complexity.
  - `objective='multi:softprob'` : Specifies multi-class classification with probability outputs.
- **Training Process:** The dataset was split into an 80% training set and a 20% test set using `sklearn.model_selection.train_test_split`. The model was trained on the training set. A `LabelEncoder` was used to convert the string labels (`Insurance_Type`) into integers required by XGBoost.

```
✓ XGBoost training complete!
```

```
=====
```

### 🚀 XGBoost Model Evaluation

```
=====
```

1. Binary Accuracy (Bought vs Not Bought): 0.9040

	precision	recall	f1-score	support
0	0.71	0.60	0.65	294
1	0.93	0.96	0.94	1706
accuracy			0.90	2000
macro avg	0.82	0.78	0.80	2000
weighted avg	0.90	0.90	0.90	2000

2. Multiclass Accuracy (Specific Type): 0.6335

	precision	recall	f1-score	support
A&H	0.57	0.44	0.50	395
FSA	0.58	0.54	0.56	549
Life	0.66	0.81	0.73	762
No_Insurance	0.71	0.60	0.65	294
accuracy			0.63	2000
macro avg	0.63	0.60	0.61	2000
weighted avg	0.63	0.63	0.63	2000

Multiclass Confusion Matrix:

	A&H	FSA	Life	No_Insurance
A&H	173	96	98	28
FSA	58	299	164	28
Life	48	77	620	17
No_Insurance	22	42	55	175

```
✓ Model and LabelEncoder saved to: model.pkl
```

### 3.3.3 Model Evaluation

The trained model was evaluated on the held-out test set using multiple metrics to assess its performance comprehensively.

```
✓ Model and LabelEncoder loaded from model.pkl
```

```
📊 Predicting on 2000 test samples...
```

```
--- First 10 Predictions ---
```

	Actual_Label	Predicted_Label	Actual_Code	Predicted_Code
0	No_Insurance	No_Insurance	3	3
1	FSA	Life	1	2
2	Life	Life	2	2
3	FSA	FSA	1	1
4	A&H	A&H	0	0
5	Life	Life	2	2
6	A&H	A&H	0	0
7	No_Insurance	No_Insurance	3	3
8	Life	Life	2	2
9	No_Insurance	No_Insurance	3	3

```
📈 Performance Metrics:
```

```
Accuracy: 0.6335
```

```
Precision: 0.6290
```

```
Recall: 0.6335
```

```
F1-score: 0.6256
```

```
📋 Classification Report:
```

	precision	recall	f1-score	support
A&H	0.57	0.44	0.50	395
FSA	0.58	0.54	0.56	549
Life	0.66	0.81	0.73	762
No_Insurance	0.71	0.60	0.65	294
accuracy			0.63	2000
macro avg	0.63	0.60	0.61	2000
weighted avg	0.63	0.63	0.63	2000

## 3.4 Cloud Deployment

### 3.4.1 Aliyun OSS Setup

Alibaba Cloud's Object Storage Service (OSS) was chosen as the cloud storage solution for its durability, scalability, and ease of integration.

- **Bucket Creation:** A dedicated OSS bucket (e.g., `insurance-prediction-system-bucket`) was created to store all project artifacts.
- **Access Control:** RAM (Resource Access Management) policies were configured to control access to the bucket, ensuring secure storage.

### 3.4.2 Artifact Upload

The following key artifacts were uploaded to the OSS bucket:

1. **Dataset:** The full, processed CSV dataset (`customer_insurance_dataset.csv`).
2. **Model Artifact:** The trained XGBoost model, along with the `LabelEncoder` and feature preprocessing objects, were serialized into a single pickle file (`model.pkl`) using `joblib`. This "model pipeline" encapsulates all preprocessing steps and the model itself, enabling end-to-end predictions.
3. **Metadata:** A `README.md` file documenting the contents of the bucket, data dictionaries, and instructions for model loading and inference.

### 3.4.3 Deployment Workflow

The deployment process was automated via Python scripts using the `oss2` SDK for Alibaba Cloud:

1. **Local Validation:** The model and dataset were validated locally to ensure they functioned as expected.
2. **SDK Initialization:** The Python script initialized the OSS client using access keys.
3. **File Upload:** The script uploaded `customer_insurance_dataset.csv` and `model.pkl` to the specified OSS bucket.
4. **Post-Upload Verification:** The script verified the successful upload by checking the file existence and metadata in the OSS bucket.

## 4. System Testing and Optimization Results

### 4.1 Database Performance Optimization

The effectiveness of the database optimization strategies (indexing and partitioning) was quantified through performance testing.

### 4.2 Model Prediction Performance Analysis

Beyond the aggregate metrics, a deeper analysis of the model's predictions was conducted:

- **Key Insights:**
  - The model demonstrated highest accuracy in predicting `FSA` for high-income customers, aligning with the rule-based logic used in data generation.

- `No_Insurance` predictions were also strong for young, single customers.
- **Error Analysis:**
  - The main source of error was confusion between `Life` and `A&H` insurance, particularly for customers in the 35–50 age bracket with dependents. This suggests these two classes share similar profiles in the dataset.
  - A notable error was observed in predicting `No_Insurance` for customers aged 25–30 with high incomes but no dependents. The model sometimes incorrectly predicted `FSA` for this subgroup, indicating that the income feature might be overly influential.

## 5. Summary and Future Work

### 5.1 Project Achievements

This project successfully delivered a complete, end-to-end Customer Insurance Prediction System:

1. **A Robust Data Foundation:** A well-designed and optimized PostgreSQL database was implemented, capable of efficiently storing and querying insurance business data.
2. **A Realistic Synthetic Dataset:** A large, feature-rich dataset was constructed, emulating real customer behavior, which can be used for further model development and research.
3. **A Predictive ML Model:** An XGBoost classifier was trained and evaluated, achieving a prediction accuracy of `63%` on unseen data, demonstrating its potential to assist in insurance product recommendation.
4. **Cloud-Ready Deployment:** The model and dataset were successfully deployed to Alibaba Cloud OSS, making them accessible for integration into production applications.

### 5.2 Limitations and Future Directions

Despite the project's successes, several avenues for improvement and future exploration exist:

1. **Data Enhancement:**

- **Incorporate Real Data:** Transition from synthetic to anonymized real-world customer data to improve model generalization.
- **Feature Expansion:** Introduce more granular features such as credit score, occupation, geographic location, and past insurance claims history.

## 2. Model Improvement:

- **Advanced Algorithms:** Experiment with more complex models like Neural Networks (e.g., using TensorFlow/PyTorch) or ensemble methods combining XGBoost with other models.
- **Hyperparameter Optimization:** Utilize techniques like Grid Search or Bayesian Optimization for more systematic hyperparameter tuning.
- **Handling Class Imbalance:** If real data shows significant class imbalance, employ techniques like SMOTE (Synthetic Minority Oversampling Technique) or class weights.

## 3. System Scalability & Operationalization:

- **Real-Time Inference:** Deploy the model as a REST API endpoint (e.g., using Flask, FastAPI, or Aliyun Function Compute) to enable low-latency, real-time predictions.
- **Automated Pipeline:** Implement an MLOps pipeline using tools like Airflow or Kubeflow to automate data ingestion, preprocessing, model training, evaluation, and deployment.
- **Interactive Dashboard:** Develop a business intelligence dashboard (e.g., using Streamlit, Dash, or Power BI) to visualize model predictions, feature importance, and key business metrics for stakeholders.

## 4. Ethical AI and Bias Mitigation:

Conduct a thorough audit of the model for potential biases (e.g., gender, age) and implement fairness-aware machine learning techniques to ensure equitable predictions across all customer segments.